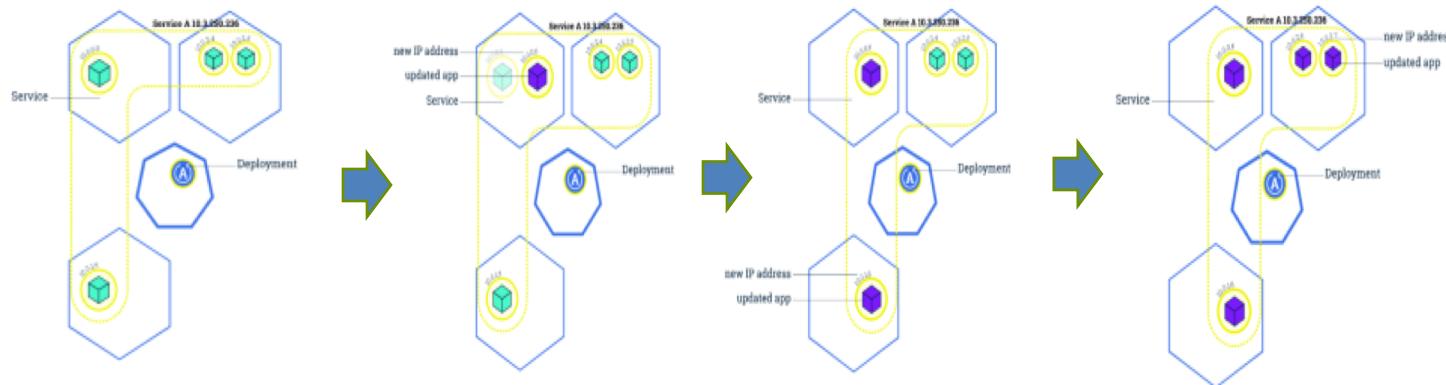


# Chapter 5: Kubernetes for Continuous Integration

# Kubernetes: Rolling Update

**Updates the pods in a serial manner will load balancing traffic to available instances**

- If deployment is exposed publicly, the service will load-balance traffic to only active and available pods
- Used to enable zero downtime updates
- Allows greater application update frequency
- Can also be leveraged to rollback to previous versions



# Zero Downtime Rolling updates and rollback

# \$ kubectl set image <deployment> <new-image>

```
peter@Azure:~$ kubectl run nginx --image nginx:1.12.1 --port 80
deployment "nginx" created
peter@Azure:~$ kubectl set image deployments/nginx nginx=nginx:latest
deployment "nginx" image updated
peter@Azure:~$ kubectl describe deployment nginx
Name:           nginx
Namespace:      default
CreationTimestamp:   Fri, 15 Sep 2017 13:08:50 +0000
Labels:          run=nginx
Annotations:    deployment.kubernetes.io/revision=2
Selector:        run=nginx
Replicas:       1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
  Labels:       run=nginx
  Containers:
    nginx:
      Image:      nginx:latest
      Port:       80/TCP
      Environment:  <none>
      Mounts:      <none>
      Volumes:     <none>
  Conditions:
    Type        Status  Reason
    ----        ----   -----
    Available   True    MinimumReplicasAvailable
OldReplicaSets: <none>
NewReplicaSet:  nginx-2688028062 (1/1 replicas created)
Events:
FirstSeen  LastSeen  Count  From            SubObjectPath  Type        Reason               Message
-----  -----  -----  -----            -----  -----  -----   -----
32s       32s       1      deployment-controller      Normal      ScalingReplicaSet  Scaled up replica set nginx-1295584306 to 1
10s       10s       1      deployment-controller      Normal      ScalingReplicaSet  Scaled up replica set nginx-2688028062 to 1
10s       10s       1      deployment-controller      Normal      ScalingReplicaSet  Scaled down replica set nginx-1295584306 to 0
```

**Note:** Rolling updates can also be undone (see next slide)

# \$ kubectl rollout undo <deployment>

```
peter@Azure:~$ kubectl rollout undo deployments/nginx
deployment "nginx" rolled back
peter@Azure:~$ kubectl describe deployments/nginx
Name:           nginx
Namespace:      default
CreationTimestamp:   Fri, 15 Sep 2017 13:36:49 +0000
Labels:          run=nginx
Annotations:    deployment.kubernetes.io/revision=3
Selector:        run=nginx
Replicas:       1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
  Labels:       run=nginx
  Containers:
    nginx:
      Image:      nginx:1.12.1
      Port:       80/TCP
      Environment:  <none>
      Mounts:     <none>
      Volumes:    <none>
  Conditions:
    Type        Status  Reason
    ----        ----   -----
    Available   True    MinimumReplicasAvailable
OldReplicaSets: <none>
NewReplicaSet:  nginx-1295584306 (1/1 replicas created)
Events:
FirstSeen  LastSeen  Count  From            SubObjectPath  Type        Reason               Message
-----  -----  -----  -----            -----  -----  -----  -----
1m         1m        1  deployment-controller  Normal      ScalingReplicaSet  Scaled up replica set nginx-2688028062 to 1
1m         1m        1  deployment-controller  Normal      ScalingReplicaSet  Scaled down replica set nginx-1295584306 to 0
1m         10s       2  deployment-controller  Normal      ScalingReplicaSet  Scaled up replica set nginx-1295584306 to 1
10s        10s       1  deployment-controller  Normal      DeploymentRollback  Rolled back deployment "nginx" to revision 1
10s        10s       1  deployment-controller  Normal      ScalingReplicaSet  Scaled down replica set nginx-2688028062 to 0
```

**Note:** Rollbacks can also be also enacted with zero-downtime

```

root@ip-172-31-0-112:~/code# kubectl run nginx --image nginx:1.12.1 --port 80
deployment.apps/nginx created
root@ip-172-31-0-112:~/code# kubectl set image deployments/nginx nginx=nginx:latest
deployment.extensions/nginx image updated
root@ip-172-31-0-112:~/code# kubectl describe deployment nginx
Name:           nginx
Namespace:      default
CreationTimestamp: Thu, 19 Jul 2018 04:00:55 +0000
Labels:          run=nginx
Annotations:    deployment.kubernetes.io/revision=2
Selector:        run=nginx
Replicas:       1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  run=nginx
  Containers:
    nginx:
      Image:      nginx:latest
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Conditions:
    Type     Status  Reason
    ----     ----   -----
    Available  True    MinimumReplicasAvailable
    Progressing True    NewReplicaSetAvailable
  OldReplicaSets: <none>
  NewReplicaSet:  nginx-6c486b77db (1/1 replicas created)
Events:
  Type     Reason            Age   From                  Message
  ----     ----            ----  ----                  -----
  Normal   ScalingReplicaSet  43s   deployment-controller  Scaled up replica set nginx-7f9bc86464 to 1
  Normal   ScalingReplicaSet  11s   deployment-controller  Scaled up replica set nginx-6c486b77db to 1
  Normal   ScalingReplicaSet  10s   deployment-controller  Scaled down replica set nginx-7f9bc86464 to 0

```



```

root@ip-172-31-0-112:~/code# kubectl rollout undo deployments/nginx
deployment.extensions/nginx
root@ip-172-31-0-112:~/code# kubectl describe deployments/nginx
Name:           nginx
Namespace:      default
CreationTimestamp: Thu, 19 Jul 2018 04:00:55 +0000
Labels:          run=nginx
Annotations:    deployment.kubernetes.io/revision=3
Selector:        run=nginx
Replicas:       1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  run=nginx
  Containers:
    nginx:
      Image:      nginx:1.12.1
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Conditions:
    Type      Status  Reason
    ----      ----   -----
    Available  True    MinimumReplicasAvailable
    Progressing True    NewReplicaSetAvailable
    OldReplicaSets: <none>
    NewReplicaSet:  nginx-7f9bc86464 (1/1 replicas created)
  Events:
    Type      Reason     Age           From            Message
    ----      ----     --           ----           -----
    Normal   ScalingReplicaSet  1m          deployment-controller  Scaled up replica set nginx-6c486b77db to 1
    Normal   ScalingReplicaSet  1m          deployment-controller  Scaled down replica set nginx-7f9bc86464 to 0
    Normal   ScalingReplicaSet  18s (x2 over 2m)  deployment-controller  Scaled up replica set nginx-7f9bc86464 to 1
    Normal   DeploymentRollback 18s          deployment-controller  Rolled back deployment "nginx" to revision 1
    Normal   ScalingReplicaSet  17s          deployment-controller  Scaled down replica set nginx-6c486b77db to 0

```

**Note:** Rollbacks can also be enacted with zero-downtime

kubectl rollout status deployment nginx

kubectl rollout history deployment nginx

kubectl rollout history deployment/nginx --revision=3

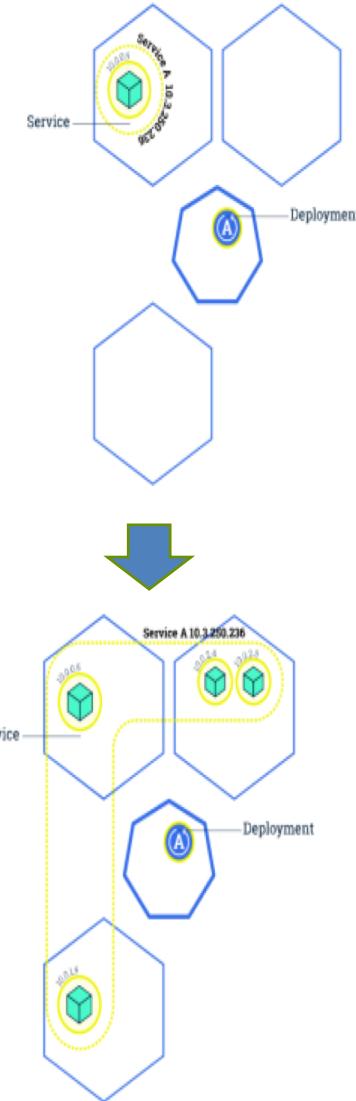
kubectl rollout undo deployment/nginx --to-revision <num>

# Zero Downtime Scaling

# Kubernetes: Scaling

**Changing the number of resources to meet a desired state**

- Accomplished by adjusting the number of replicas in a deployment
- Accommodates both scaling up and scaling down to a minimum of 0
- Traffic is automatically sent to newly created instances through the service load balancer
- Can be used to enable rolling updates without downtime



# \$ kubectl scale deployments/<deployment> --replicas=<new num>

\$ kubectl scale deploy/nginx --replicas=4

```
peter@Azure:~$ kubectl run nginx --image nginx --port 80
deployment "nginx" created
peter@Azure:~$ 
peter@Azure:~$ kubectl get deployments
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx     1          1          1           1           39s
peter@Azure:~$ kubectl scale deployments/nginx --replicas=2
deployment "nginx" scaled
peter@Azure:~$ kubectl get deployments
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx     2          2          2           1           1m
```

Note: Make sure to delete the deployments or replicaset when trying to clear out pods. Many a new user has repeated deleted pods and gotten frustrated when they respawn (as the deployments/replicaset are programed to do).

```

root@ip-172-31-0-112:~/code# kubectl run nginx --image nginx --port 80
deployment.apps/nginx created
root@ip-172-31-0-112:~/code#
root@ip-172-31-0-112:~/code# kubectl get deployments
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx     1          1          1           1           8s
nginx-deployment 3          3          3           3           17m
redis     1          1          1           1           33m
root@ip-172-31-0-112:~/code#
root@ip-172-31-0-112:~/code# kubectl scale deployments/nginx --replicas=3
deployment.extensions/nginx scaled
root@ip-172-31-0-112:~/code# kubectl get deployments
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx     3          3          3           3           34s
nginx-deployment 3          3          3           3           18m
redis     1          1          1           1           33m
root@ip-172-31-0-112:~/code# kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
nginx-6f858d4d45-b4cv2        1/1    Running   0          42s
nginx-6f858d4d45-g42mg        1/1    Running   0          16s
nginx-6f858d4d45-169jv        1/1    Running   0          16s
nginx-apparmor                 1/1    Running   0          30m
nginx-deployment-67594d6bf6-7v5lk 1/1    Running   0          12m
nginx-deployment-67594d6bf6-fdqx2 1/1    Running   0          18m
nginx-deployment-67594d6bf6-sx62p 1/1    Running   0          18m
redis-7869f8966-sq8xm         1/1    Running   0          33m

```

# Autoscaling and HPA (Horizontal Pod Autoscaler)

# Kubernetes: AutoScaling

## HPA – Horizontal Pod Autoscaler:

Based on memory and CPU utilization by a pod, autoscaling scales up or down the number of pods

```
$ kubectl autoscale deployment nginx --min=2 --max=10
```

```
$ kubectl autoscale deployment nginx-deployment --max=5 --cpu-percent=80
```

```
$ kubectl get hpa
```

```
$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
nginx	Deployment/nginx	<unknown>/80%	1	4	3	12m
nginx-deployment	Deployment/nginx-deployment	<unknown>/80%	2	5	0	5s

# Exercise 4 – (scaling, Auto scaling, rollout)

- For a previously created deployment scale the number of replicas
- see the difference in replicsets, pods , and see with a describe
- Rolling update, undo and see the changes in replica sets

cd to the ex4

[https://github.com/shekhar2010us/kubernetes\\_teach\\_git/blob/master/ex4/scaling\\_rolling.md](https://github.com/shekhar2010us/kubernetes_teach_git/blob/master/ex4/scaling_rolling.md)

# Service Discovery

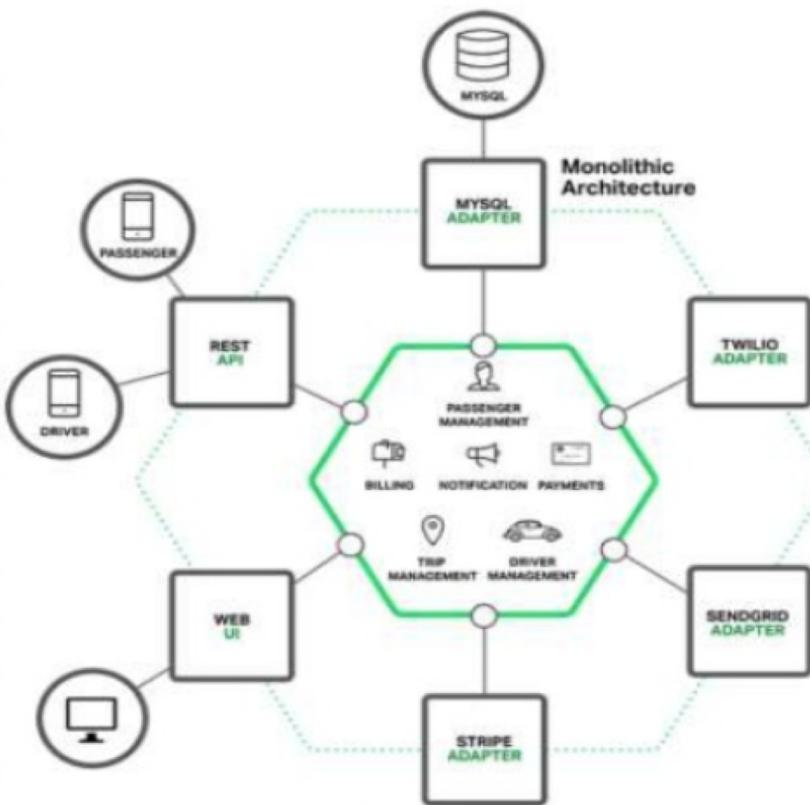
# Service Discovery

Migrating from a monolithic architecture to a Microservices based architecture where each service loosely coupled and horizontally scaled will require a service discovery mechanism to ensure that the REST API endpoints of the ephemeral services (say, docker containers running inside PODs) are resolved or detected. The IP addresses of the containers (or the PODs) being transient will need automated mechanism to resolve or convey the current IP addresses.

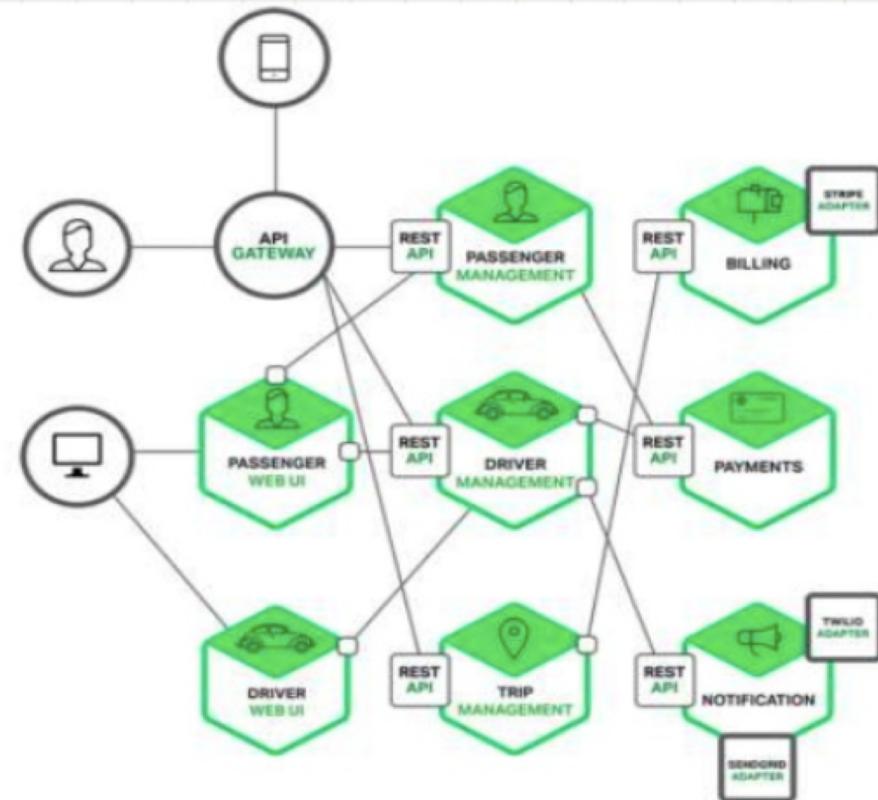


# Sample Microservices Architecture

Monolith



Microservices



# Service Discovery

**So how do we keep  
track of all the  
microservices?**



**Service Discovery!**



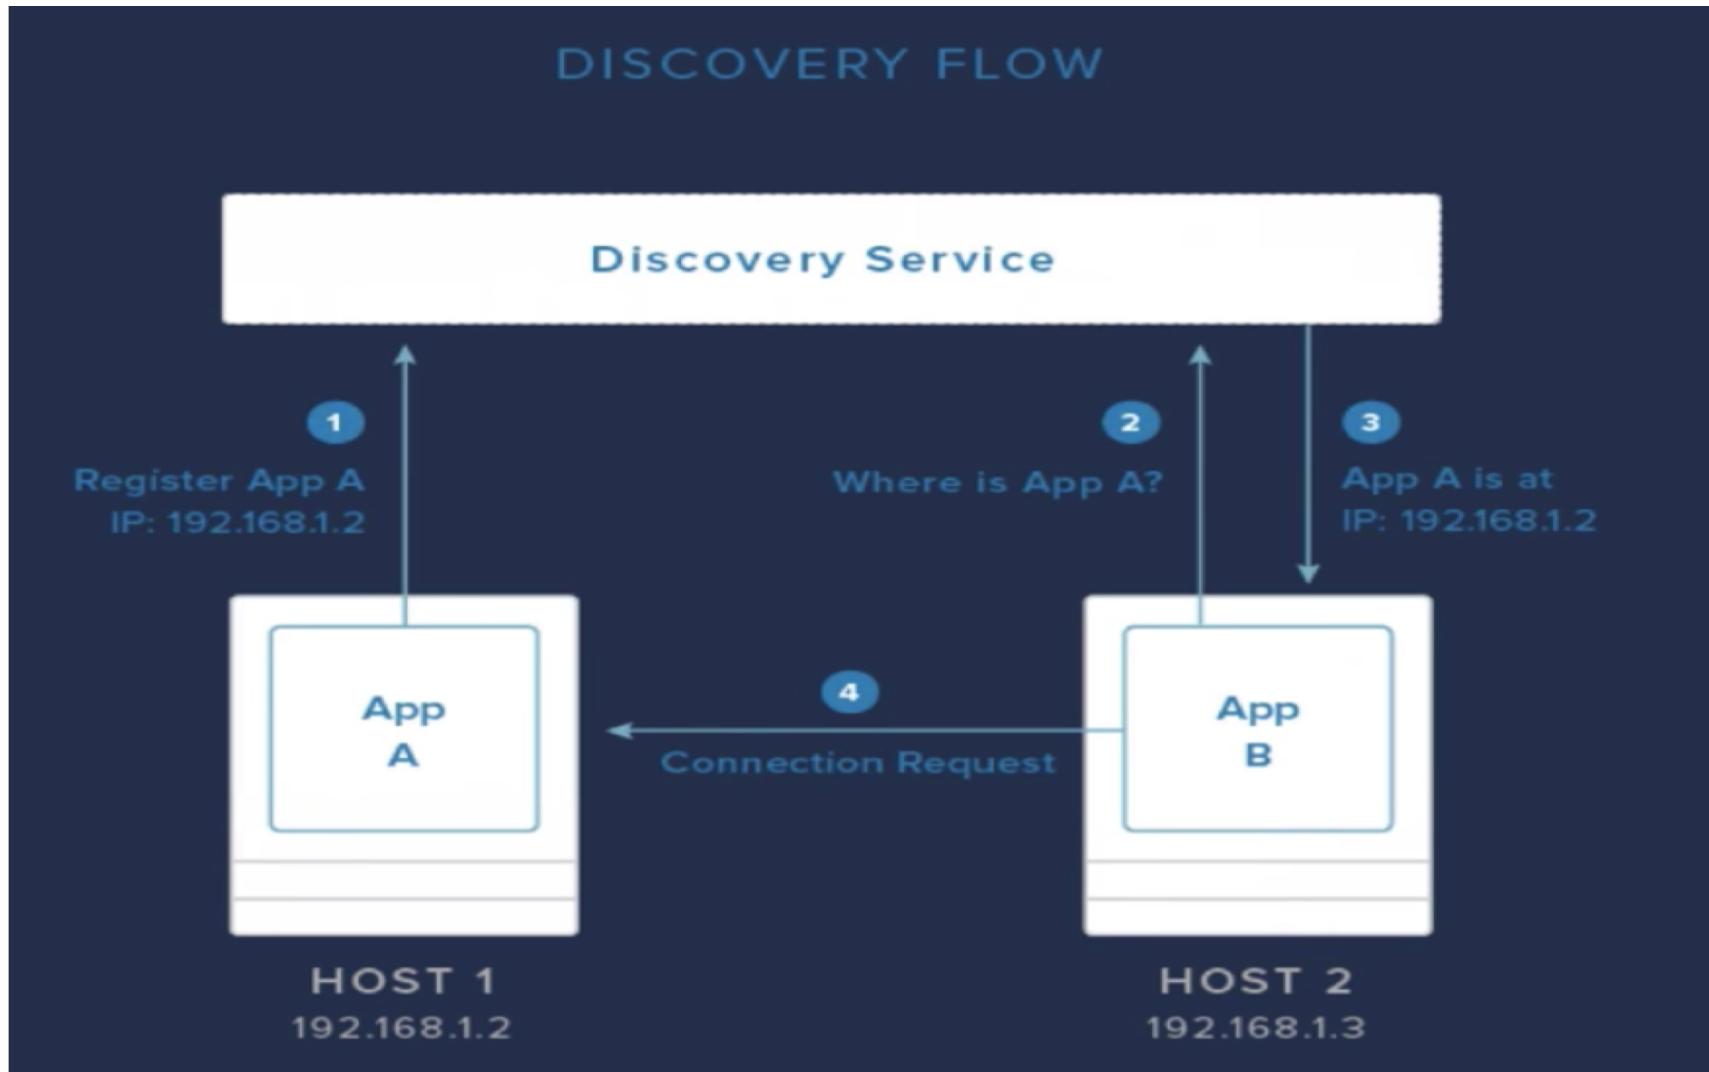
# Service Discovery

What should Service Discovery provide?

- **Discover services** dynamically to get IP address and port details to communicate with other services
- **Health check:** Only healthy services should participate in handling traffic
- **Load balancing:** Traffic destined to a particular service should be dynamically load balanced



# Service Discovery



# Service Discovery Tools

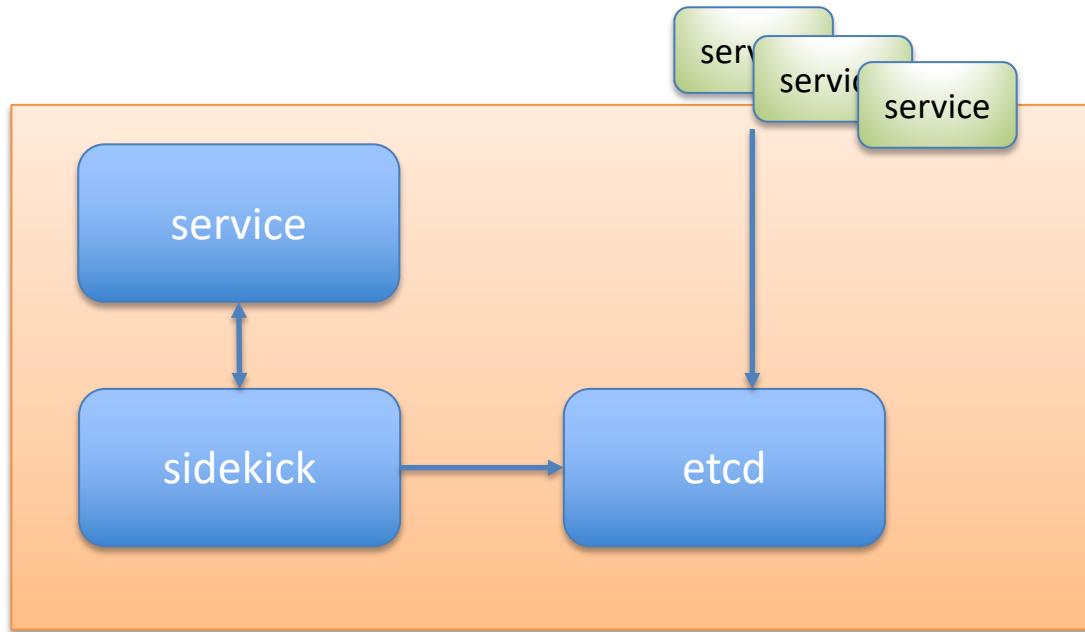
## What are some Service Discovery tools?

- etcd, by CoreOS
- Consul, by Hashicorp
- Zookeeper, by Apache
- SkyDNS (built on top of etcd)
- Eureka, by Netflix
- Smartstack, by AirBnB



# Service Discovery by etcd

- Etcd can be used as KV store for service registry
- Service itself can directly update etcd or a Sidekick service can be used to update etcd
- Sidekick service serves as **registrator**
- Other services can query etcd database to do the dynamic service discovery



# Service Discovery by HashiCorp (Consul)

- Consul is a distributed, HA service discovery and configuration system
- Consul provides features - includes service and node discovery mechanisms, tagging system, health checks, consensus-based election routines, system-wide key/value storage, etc.
- Provides Rest based HTTP Api for external interaction

# Service Discovery – Zookeeper by Apache

- Apache ZooKeeper is a distributed, open-source coordination service for distributed applications.
- Allows to read, write, and observe updates to data.
- Data are organized in a file system like hierarchy and replicated to all ZooKeeper servers in the ensemble (a set of ZooKeeper servers).
- Operations on data are atomic and sequentially consistent

# Kubernetes Service Discovery

Using the environment variables that use the same conventions as those created by Docker links. Using DNS to resolve the service names to the service's IP address.

## 1. Environment Variables

Kubernetes injects environment variables for each service and each port exposed by the service. This makes it easy to deploy containers that use Docker links to find their dependencies.

For example, a RabbitMQ service can be located by using the

**RABBITMQ\_SERVICE\_SERVICE\_HOST** and **RABBIT\_MP\_SERVICE\_SERVICE\_PORT** variables.

Other environment variables are also exposed to support this.

# Kubernetes Service Discovery

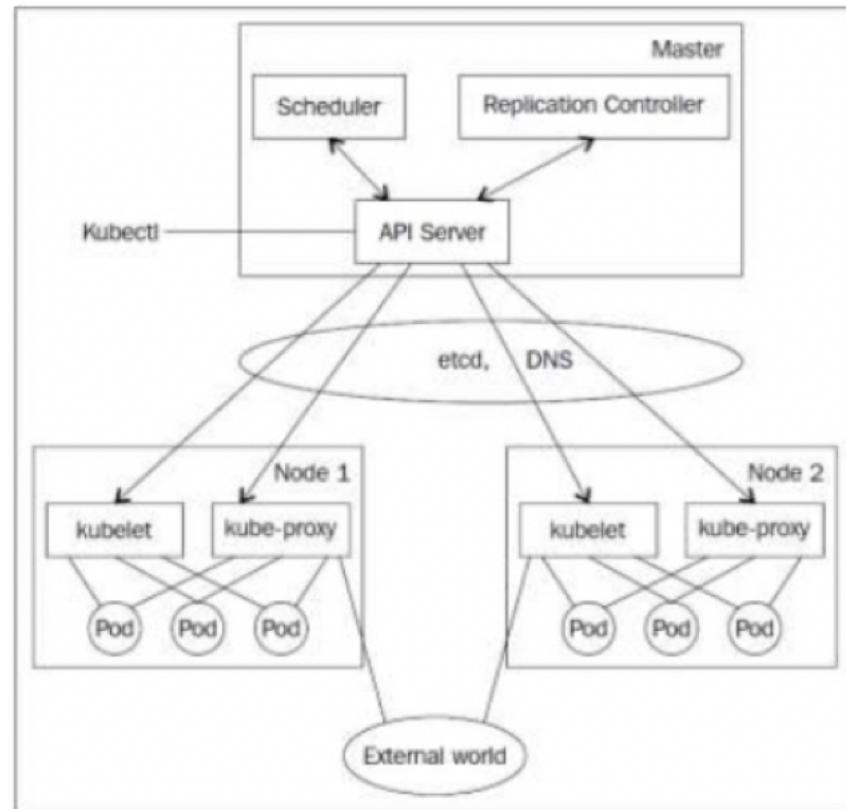
## 2. DNS Resolution

Kubernetes has a kube-dns addon that exposes the service's name as a DNS entry. As a result, you can tell your application to connect to a host name.

The service names are scoped within namespaces. This allows you to run different deployment of a service for each namespace (for example, one per developer or one per environments) without having to edit configuration files.

# K8s Service Discovery Components

- **SkyDNS** – map Service name to IP address
- **Etcd** – KV store for service database
- **Kubelet** – healthcheck and replication controller takes care of maintaining pod count
- **Kube-proxy** – takes care of load balancing traffic to individual pods.  
Watches service changes and updates IPTables



# Service Discovery

- **Services** – get mapped to pods using Selector labels. In the example: “MyApp” is the label
- **Service port** gets mapped to targetPort in the pod

```
{  
  "kind": "Service",  
  "apiVersion": "v1",  
  "metadata": { "name": "my-service" },  
  "spec": {  
    "selector":{ "app": "MyApp" },  
    "ports": [  
      {  
        "protocol": "TCP",  
        "port": 80,  
        "targetPort": 9376  
      }  
    ]  
  }  
}
```



# Lab: Service Discovery

[https://github.com/shekhar2010us/kubernetes\\_teach\\_git/blob/master/k8s\\_service\\_discovery/service-discovery.md](https://github.com/shekhar2010us/kubernetes_teach_git/blob/master/k8s_service_discovery/service-discovery.md)

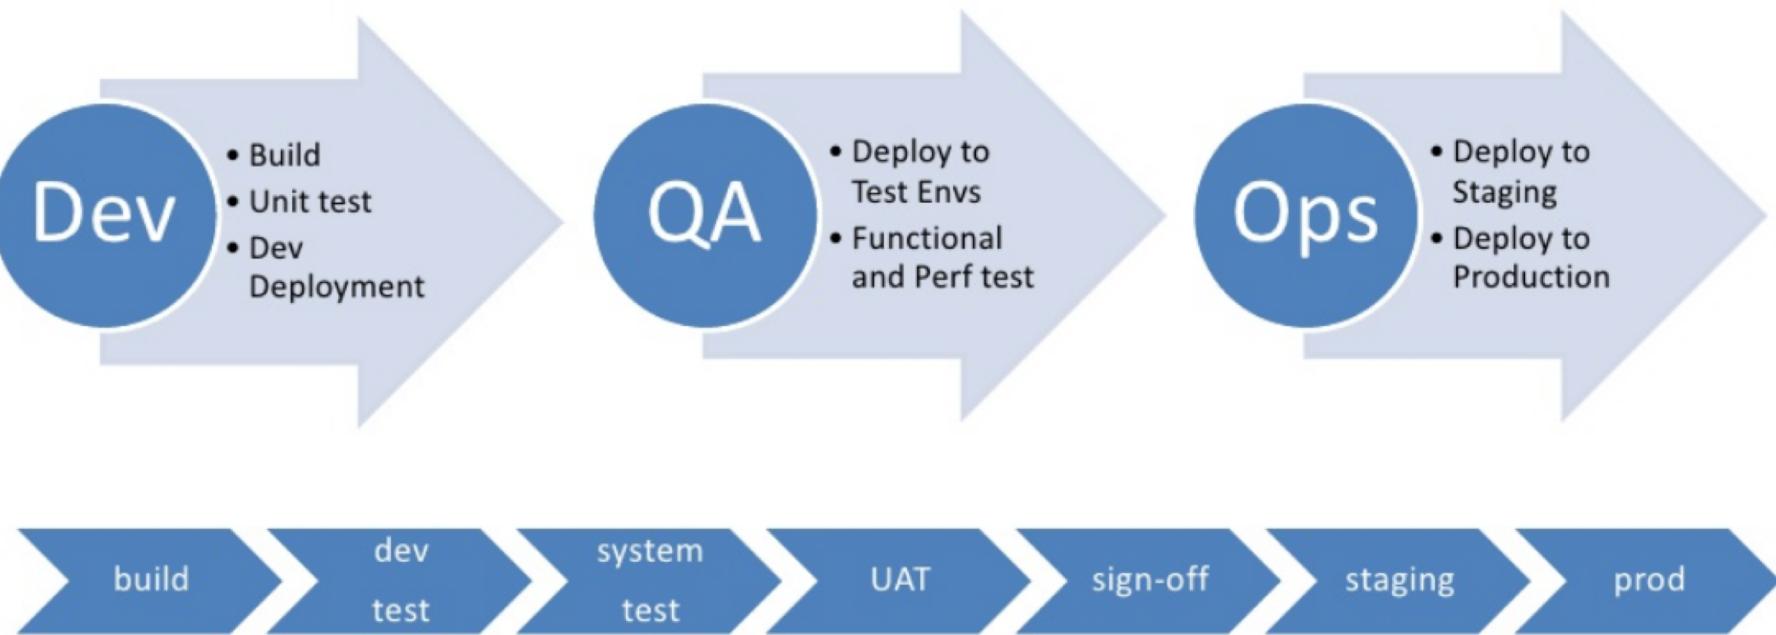
**Example using consul:**

<https://github.com/hashicorp/demo-consul-101/tree/master/k8s>

# CICD (Jenkins) Integration with Kubernetes

# Introduction to CI/CD

## What is a Pipeline



# Continuous Integration: Benefits

**Reduced integration risk.** More often than not, working on projects means multiple people are working on the separate tasks or parts of the code. The more people, the riskier the integration. Depending on how bad the problem really is, debugging and solving the issue can be really painful and can potentially mean a lot of changes to the code. Integrating on a daily basis or even more frequently can help reduce these kinds of problems to a minimum.

**Higher code quality.** Not having to worry about the problems, and focusing more on the functionality of the code results in a higher quality product.

**The code in version control works.** If you commit something that breaks the build, you and your team get the notice immediately and the problem is fixed before anyone else pulls the “broken” code.

**Less time deploying.** Deploying projects can be very tedious and time-consuming, and automating that process makes perfect sense



# Continuous Integration: Requirements

**VCS:** The first requirement is having the version control system (VCS). There is no way around it and there shouldn't be a way around it. VCS provides a reliable method to centralize and preserve changes made to your project over time. (Git)

**CI Server:** Continuous integration server (aka build server, aka CI server) is a software tool that centralizes all your CI operations and provides a reliable and stable environment for you to build your projects on.(Jenkins)



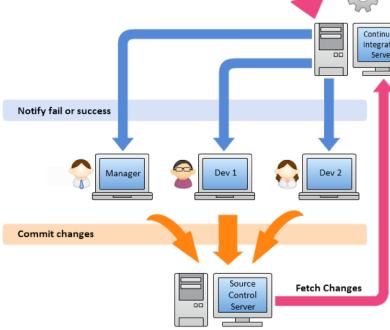
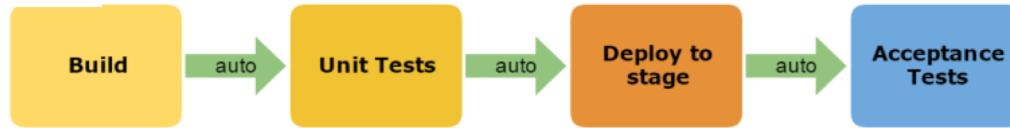
# Jenkins

Jenkins is an open source automation server written in Java. Jenkins helps to automate the non-human part of software development process, with continuous integration and continuous delivery/continuous deployment.



# Jenkins

## Continuous Integration

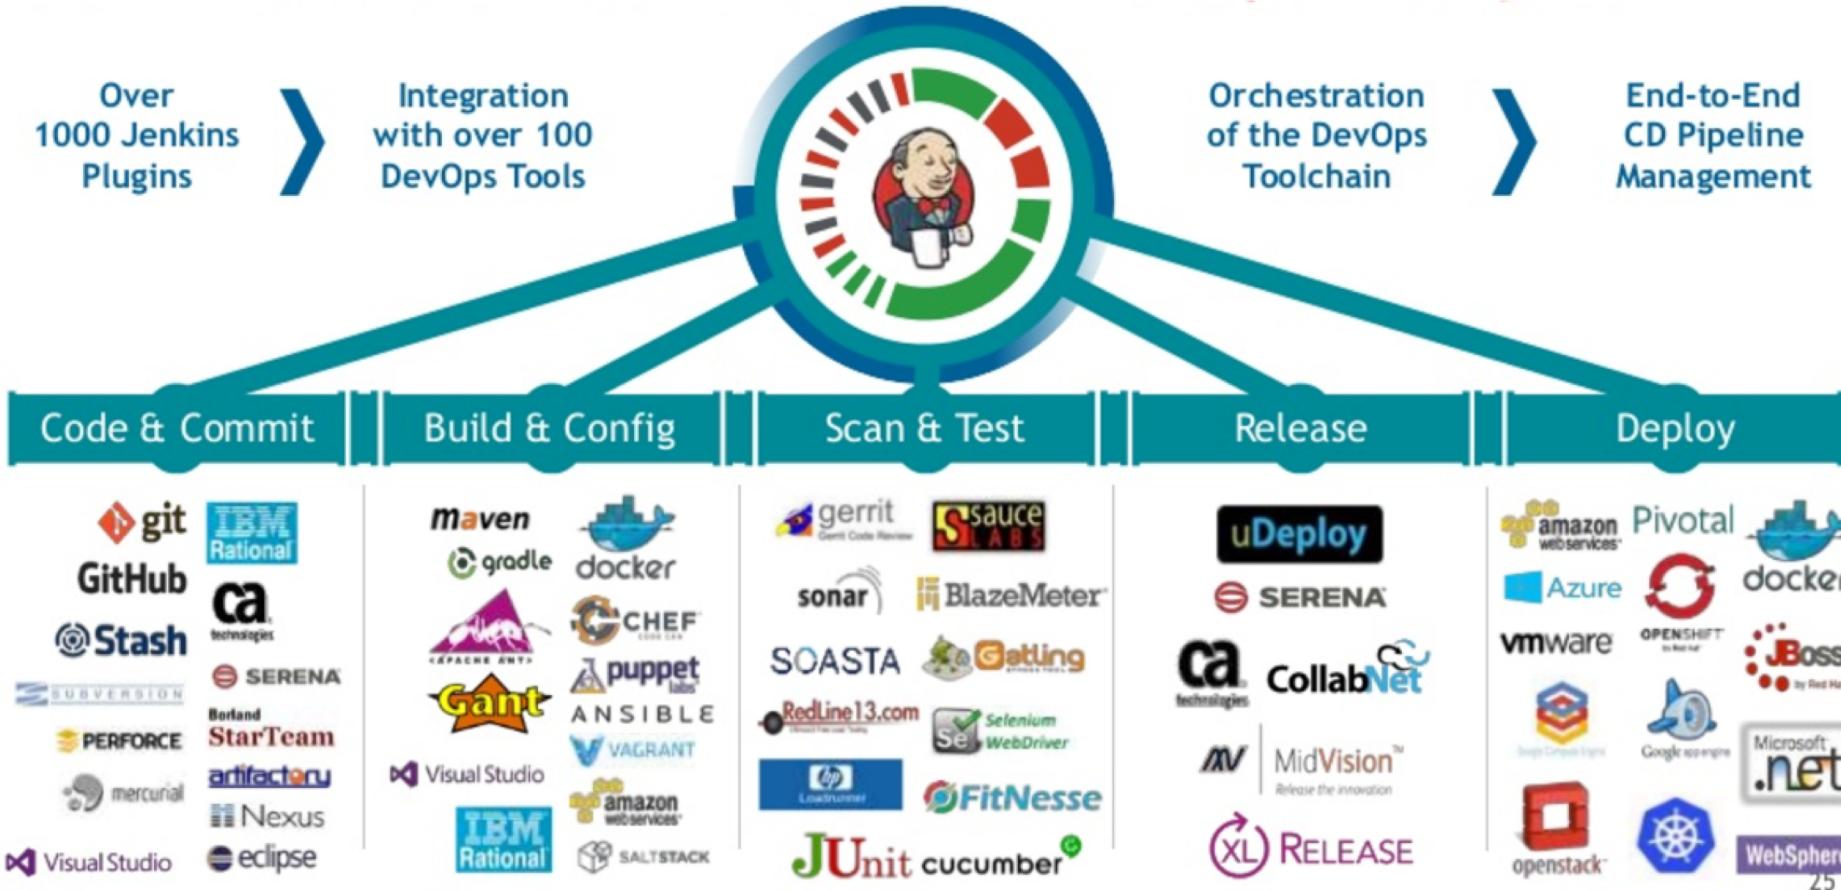


## Continuous Delivery



## Continuous Deployment





# Deploy Jenkins in Kubernetes

<https://github.com/jenkinsci/kubernetes-cd-plugin/blob/dev/README.md>

<https://medium.com/containerum/configuring-ci-cd-on-kubernetes-with-jenkins-89eab7234270>

<https://kubernetes.io/blog/2018/04/30/zero-downtime-deployment-kubernetes-jenkins/>